Self-Supervised Learning

Tutorial 19-08-2021

Supervised Learning

- The initial boost in the machine learning world came via the paradigm of supervised learning.

- In this setting, a model is trained for a specialized task for which the data is carefully labelled.
  - Bounding boxes for localization
  - Semantic maps for semantic segmentation, etc.

- Practically speaking, it's impossible to label everything in the world.

- Unfortunately, this a limits how far the field of AI can go with supervised learning alone.

Basics of Self-Supervised Learning

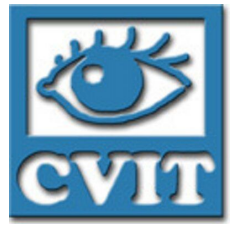

- Data provides the supervision directly.
- In general, perturb the data and task a network to predict it back.
- We often may solve a proxy task using the network forcing it to learn meaningful semantics that can be used in downstream tasks.
- The proxy tasks are also often of great research importance in standalone form.
- Let us check some image level self-supervised learning tasks.

An Example: Image Colorization

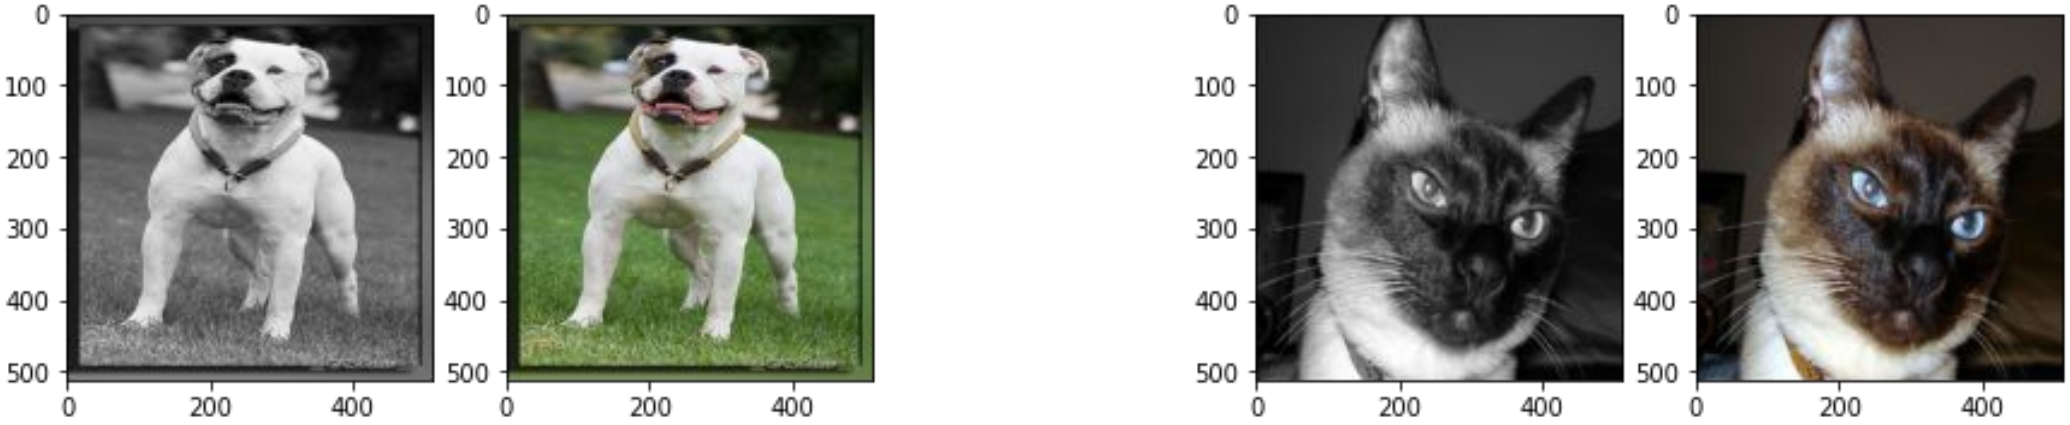

- Train a neural network to predict colours from a grayscale image.
- The network needs to inherently learn semantic boundaries present in the image, for example the shape of the foreground (dog), the background type etc.
- This semantic knowledge can be exploited in downstream tasks like semantic segmentation, a task which now needs a human to annotate every pixel present in an image.

An Example: Image Inpainting

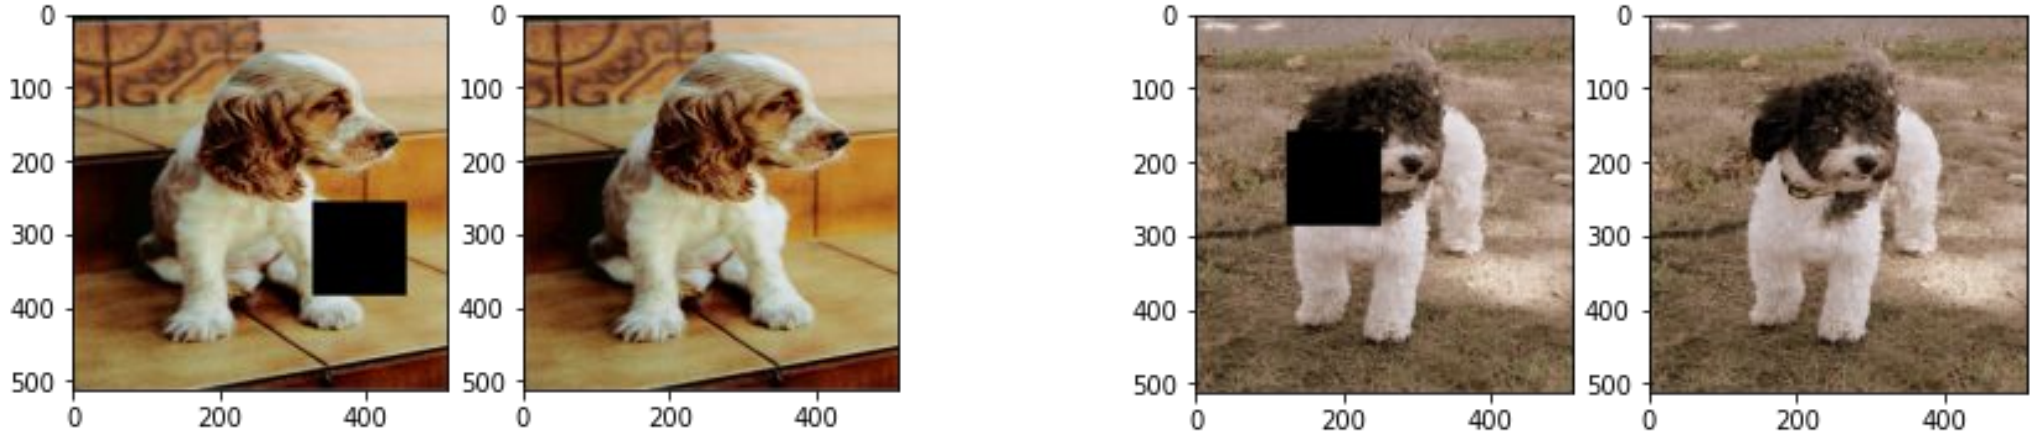

- Remove a particular area of an image randomly and ask a NN to predict it back.
- The network requires to understand the structure of objects present in the image to inpaint the required region

An Example: Image Super-resolution

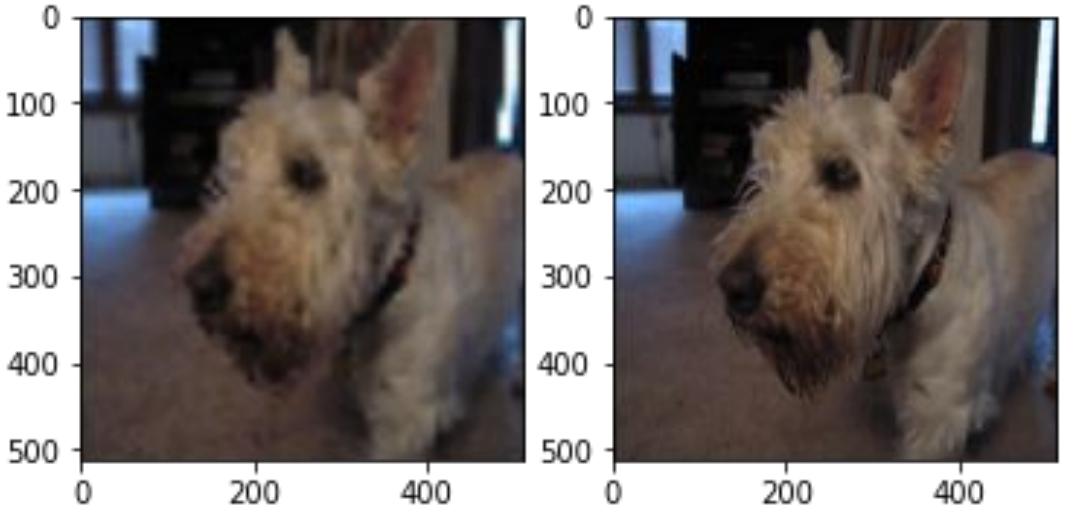

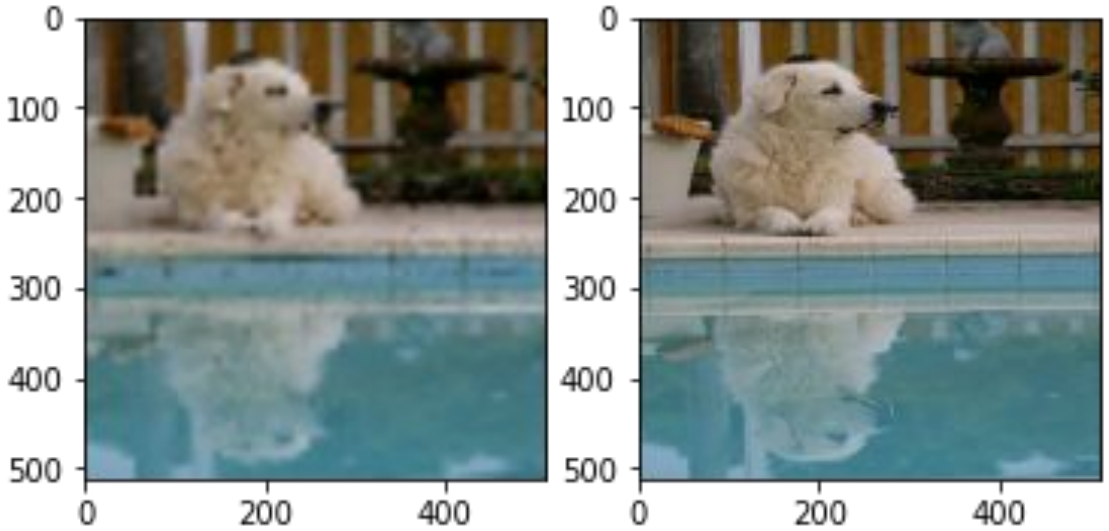

- Predicting a higher resolution image from a lower input.
- We will be showing a code walk through for this particular topic in today's session.

How are the networks trained?

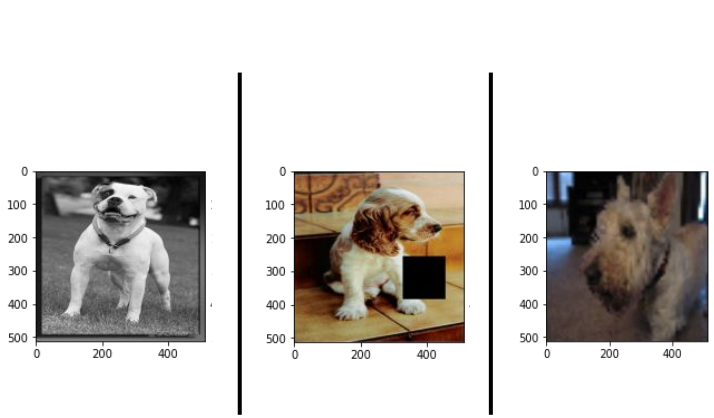

Perturbed Inputs

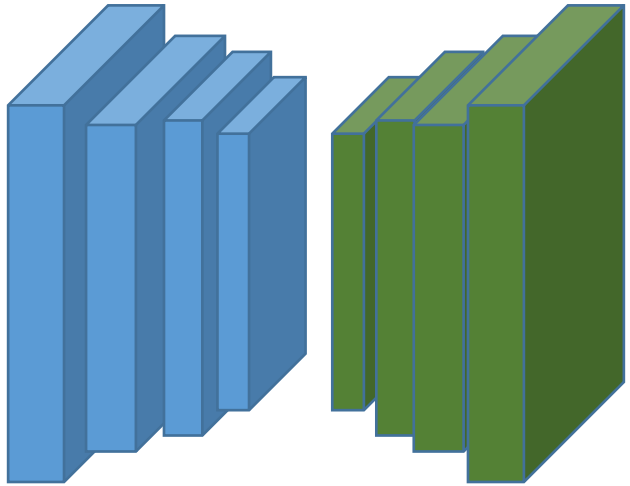

Convolutional Layers

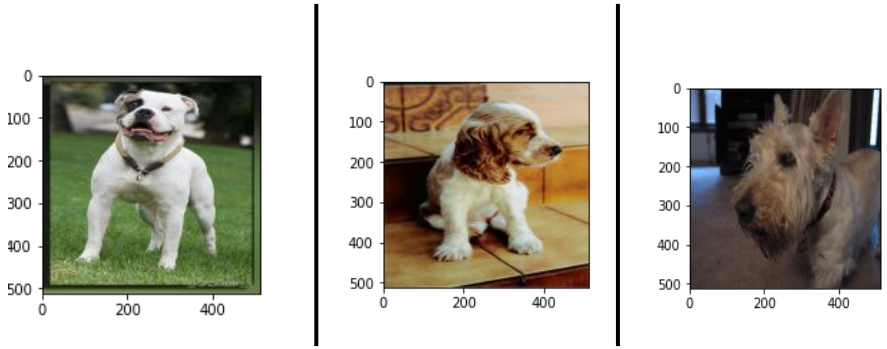

Ground Truths

Corpus of unlabelled images

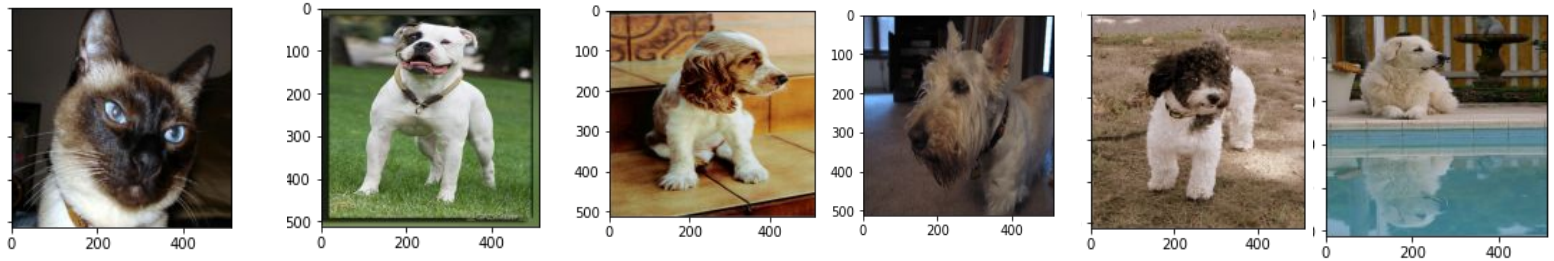

Similar approach in other fields:

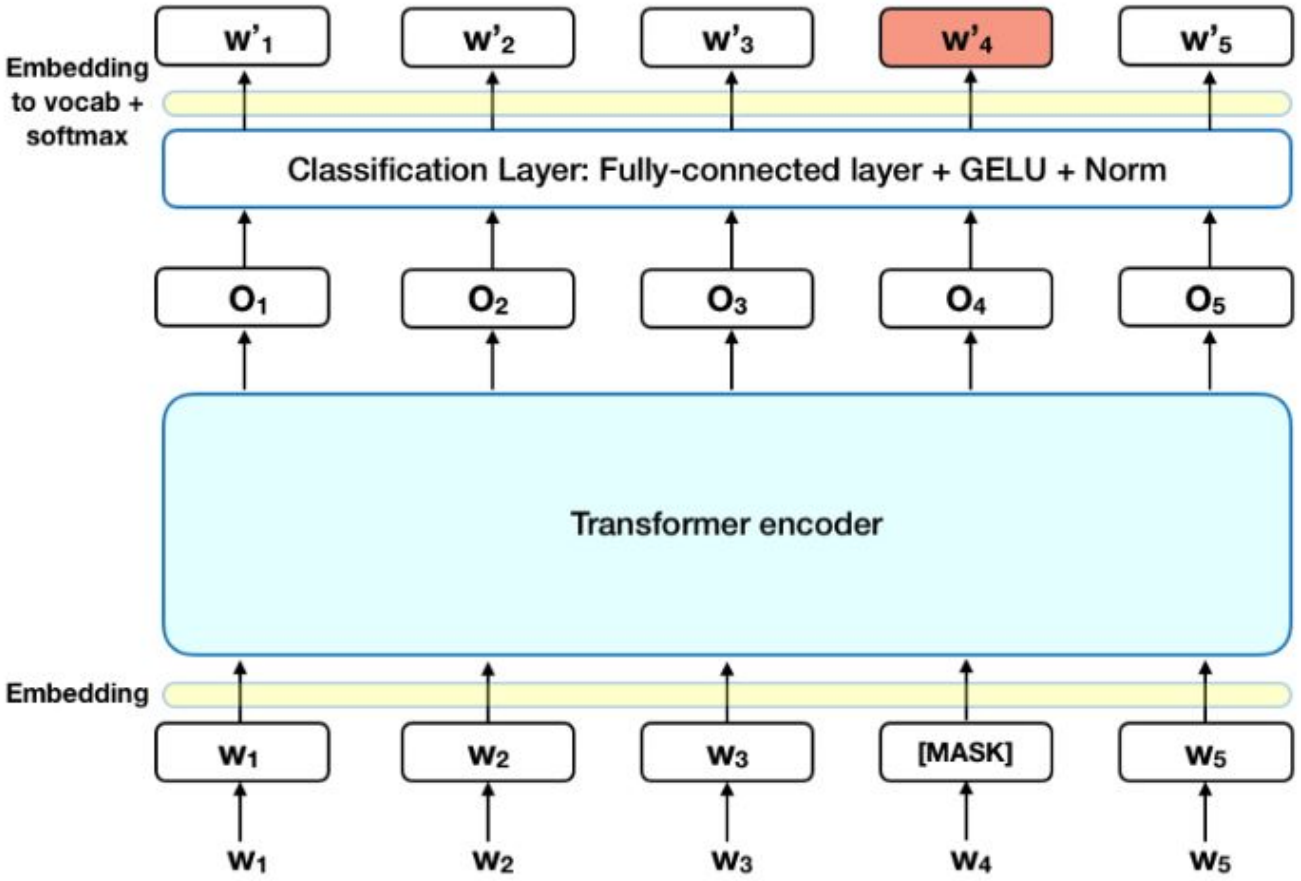

- The BERT language model is also trained on a similar concept.
- Instead of image patch, we mask a word from a random sentence.
- A transformer based network is tasked to predict the masked word back learning rich semantics present in language.
- Wav2Vec also follows a similar principal for speech.

Image credit: https://towardsdatascience.com/bert-explained-state-of-the-art-language-model-for-nlp-f8b21a9b6270

Audio-Visual Self Supervised Learning

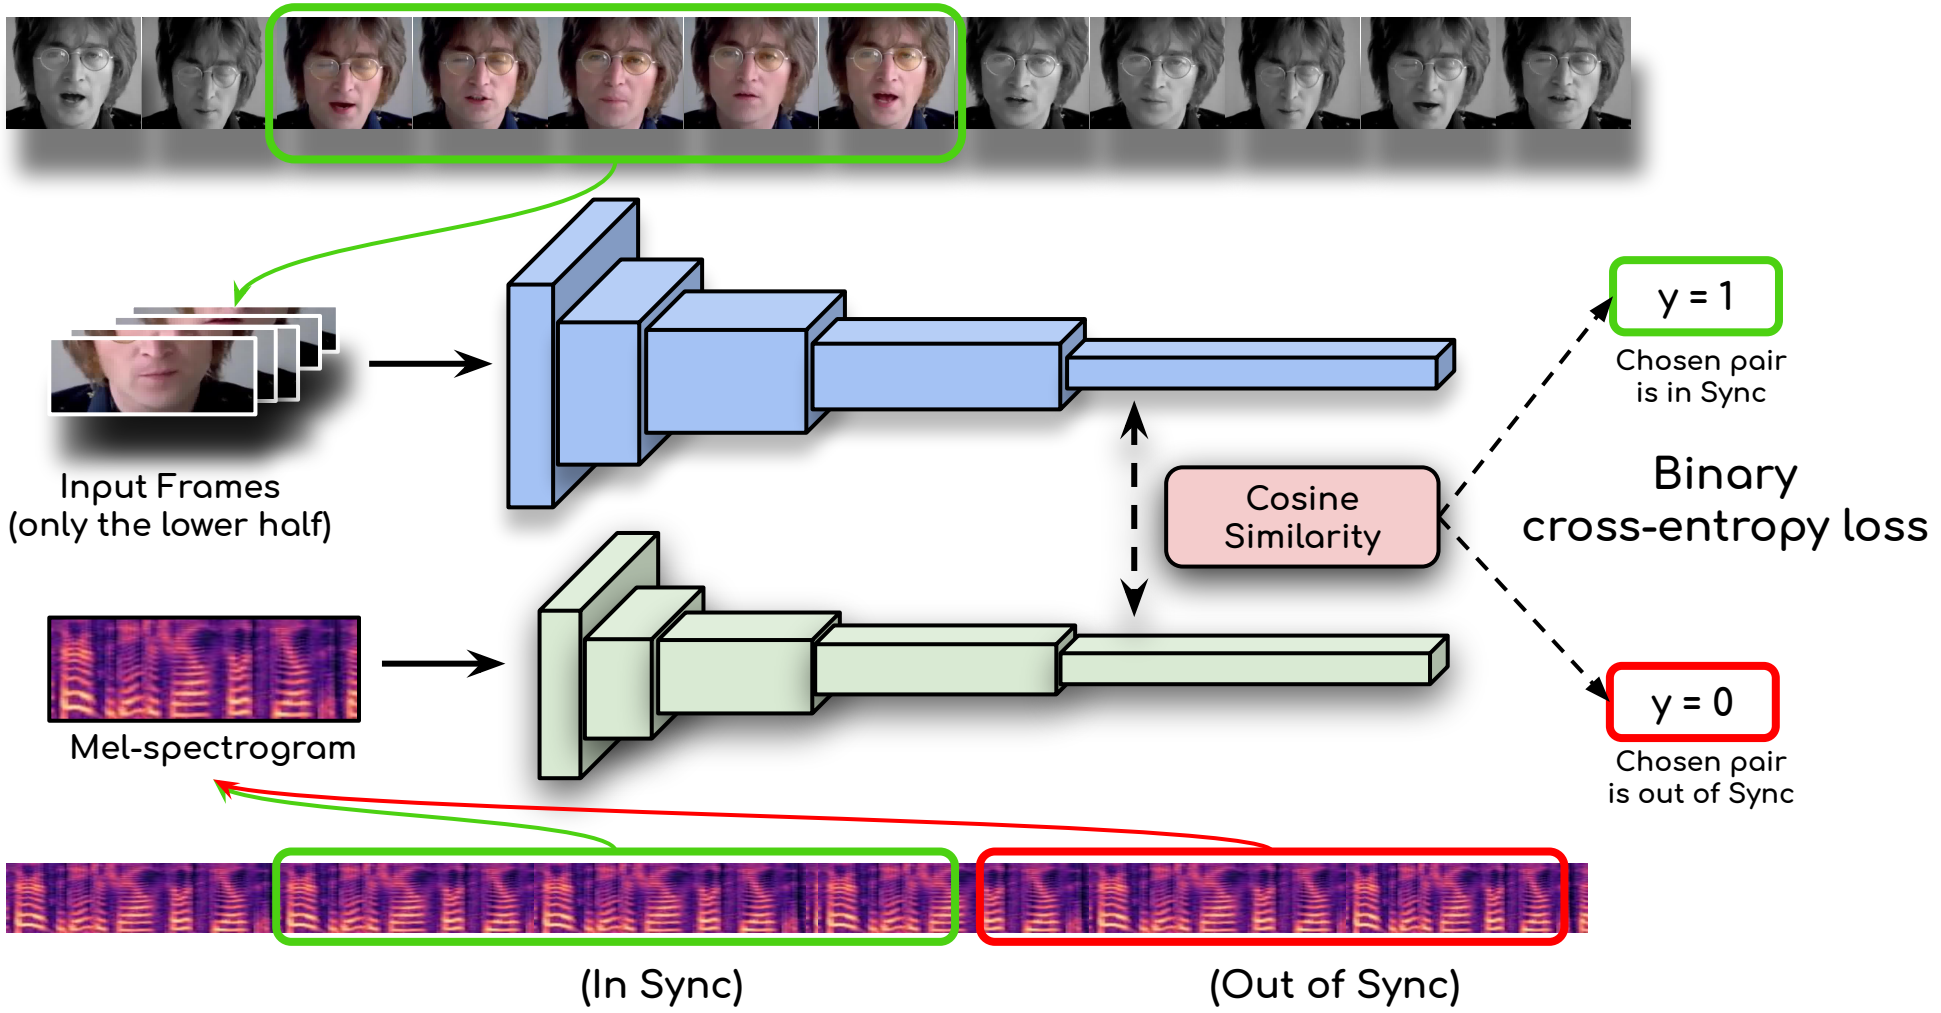

Let us now go through a SR code:

- Please go to this repository: https://github.com/Rudrabha/SS2021-19-08-2021
- Open this notebook for the code walk through: https://github.com/Rudrabha/SS2021-19-08-2021/blob/main/Image_Super_Resolve_Tutorial.ipynb
- There are other two notebooks containing codes for Image inpainting and Image Colorization.
- Please note that these codes are for basic introduction and not meant for State-of-the-art uses in any of these problems. However, the building blocks of the network can be used to train much more complex networks.
- Please check Prof. Andrew Zisserman's slides: https://project.inria.fr/paiss/files/2018/07/zisserman-self-supervised.pdf for more insights.

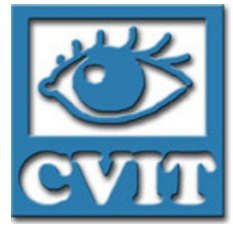
CVIT


INTERNATIONAL INSTITUTE OF
INFORMATION TECHNOLOGY
HYDERABAD

Thank You!